

OnlineCA.jl: A Julia Package for Out-of-core and Sparse Correspondence Analysis

Koki Tsuyuzaki^{1, 2}

¹ Department of Artificial Intelligence Medicine, Graduate School of Medicine, Chiba University, Japan ² Laboratory for Bioinformatics Research, RIKEN Center for Biosystems Dynamics Research, Japan

DOI: [DOIunavailable](#)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Pending Editor](#) 

Reviewers:

- [@Pending Reviewers](#)

Submitted: N/A

Published: N/A

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Correspondence Analysis (CA) is a multivariate technique for the exploratory analysis of contingency tables (Greenacre, 2017; Hill, 1974). Like Principal Component Analysis (PCA), CA decomposes a data matrix into low-dimensional row and column scores, but it operates on a standardized residual matrix that compares observed counts with the expectation under independence. CA has been widely applied across diverse fields including ecology (Braak, 1986), linguistics and text analysis (Greenacre, 2017), market research (Greenacre, 2017), and more recently single-cell omics (Hsu & Culhane, 2023).

Despite its broad applicability, CA becomes computationally prohibitive for large data matrices, since it requires the singular value decomposition (SVD) of an $N \times M$ matrix. In particular, recent advances in single-cell omics have led to datasets with millions of cells, for which standard CA implementations often fail to scale. To meet this requirement, I present `OnlineCA.jl`, a Julia package for out-of-core and sparse CA (<https://github.com/chiba-ai-med/OnlineCA.jl>).

Statement of need

CA is a workhorse algorithm for exploring contingency tables. However, as the size of the data matrix increases, it often becomes too large to fit into memory. In such cases, an out-of-core (OOC) implementation — where only subsets of data stored on disk are loaded into memory for computation — is desirable. Additionally, representing the data in a sparse matrix format, where only non-zero values and their coordinates are stored, is computationally advantageous. Therefore, a CA implementation that supports both OOC computation and sparse data handling is highly desirable (Figure 1).

Similar discussions have been made in the context of PCA and Non-negative Matrix Factorization (NMF), and we have independently developed Julia packages `OnlinePCA.jl` (Tsuyuzaki et al., 2020; Tsuyuzaki, 2026) and `OnlineNMF.jl` (Tsuyuzaki, 2025). `OnlineCA.jl` is a spin-off version of these, implementing CA.

The standardized residual matrix for CA is

$$S = D_r^{-1/2} (P - r_p c_p^\top) D_c^{-1/2},$$

where $P = X/n$ is the relative-frequency matrix, $r_p = P\mathbf{1}$ and $c_p = P^\top\mathbf{1}$ are the row- and column-mass vectors, and $D_r = \text{diag}(r_p)$, $D_c = \text{diag}(c_p)$ are the corresponding diagonal mass matrices. `OnlineCA.jl` never explicitly forms S . Instead, the matrix–vector products Sv and $S^\top u$ are computed by streaming through chunks of X on disk, and the top singular triplets are extracted with a Halko-style randomized SVD (Halko, Martinsson, & Tropp, 2011;

Halko, Martinsson, Shkolnisky, et al., 2011). The same implicit operator works uniformly for dense (CSV), sparse (Matrix Market), and binary (Binary COO) inputs.




Figure 1: Overview of workflow in OnlineCA.jl.

Example

CA can be easily reproduced on any machine where Julia is pre-installed by using the following commands in the Julia REPL window:

Installation

First, install OnlineCA.jl directly from GitHub:

```
# Install OnlinePCA.jl (used for binary I/O preprocessing)
julia> Pkg.add(url="https://github.com/rikenbit/OnlinePCA.jl.git")

# Install OnlineCA.jl
julia> Pkg.add(url="https://github.com/chiba-ai-med/OnlineCA.jl.git")
```

Preprocess of CSV

Then, write a synthetic data set as a CSV file and convert it to a compressed binary format using Zstandard. Matrix Market (MM) format is also supported for sparse matrices.

```
using OnlinePCA
using OnlinePCA: write_csv
using OnlineCA
using Distributions
using DelimitedFiles
using SparseArrays
using MatrixMarket
```

```
# CSV (input is assumed to be a non-negative integer count matrix)
tmp = mktempdir()
data = Int64.(ceil.(rand(NegativeBinomial(1, 0.5), 300, 99)))
data[1:50, 1:33] .= 100*data[1:50, 1:33]
data[51:100, 34:66] .= 100*data[51:100, 34:66]
data[101:150, 67:99] .= 100*data[101:150, 67:99]
write_csv(joinpath(tmp, "Data.csv"), data)

# Matrix Market (MM)
mmwrite(joinpath(tmp, "Data.mtx"), sparse(data))

# Binarization (Zstandard)
csv2bin(csvfile=joinpath(tmp, "Data.csv"),
        binfile=joinpath(tmp, "Data.zst"))

# Sparsification (Zstandard + MM format)
mm2bin(mmfile=joinpath(tmp, "Data.mtx"),
        binfile=joinpath(tmp, "Data.mtx.zst"))
```

Plot settings

Define a helper function to visualize the row coordinates of CA using the PlotlyJS.jl package. It generates two subplots: Component-1 vs Component-2 and Component-2 vs Component-3, with color-coded groups.

```
using DataFrames
using PlotlyJS

function subplots(out_ca, group)
    F = out_ca[1]
    data_left = DataFrame(ca1=F[:,1], ca2=F[:,2], group=group)
    data_right = DataFrame(ca2=F[:,2], ca3=F[:,3], group=group)
    p_left = Plot(data_left, x=:ca1, y=:ca2, mode="markers",
                  marker_size=10, group=:group)
    p_right = Plot(data_right, x=:ca2, y=:ca3, mode="markers",
                  marker_size=10, group=:group, showlegend=false)
    for p in (p_left, p_right)
        p.data[1]["marker_color"] = "red"
        p.data[2]["marker_color"] = "blue"
        p.data[3]["marker_color"] = "green"
    end
    p_left.data[1]["name"] = "group1"
    p_left.data[2]["name"] = "group2"
    p_left.data[3]["name"] = "group3"
    p_left.layout["title"] = "Component 1 vs Component 2"
    p_right.layout["title"] = "Component 2 vs Component 3"
    p_left.layout["xaxis_title"] = "ca-1"
    p_left.layout["yaxis_title"] = "ca-2"
    p_right.layout["xaxis_title"] = "ca-2"
    p_right.layout["yaxis_title"] = "ca-3"
    plot([p_left p_right])
end

group = vcat(repeat(["group1"], inner=100),
             repeat(["group2"], inner=100),
```

```
repeat(["group3"], inner=100))
```

CA on dense binary input

This example performs CA on the dense Zstandard-compressed binary file (Figure 2). The standardized residual matrix is never formed explicitly: `ca` streams through the data on disk to compute the implicit operator products and extracts the top `dim` components with randomized SVD.

```
out_ca = ca(input=joinpath(tmp, "Data.zst"),
            dim=3, noversamples=5, niter=3, chunksize=100)
```

```
subplots(out_ca, group)
```




Figure 2: Output of `ca` against binarized CSV format.

Sparse CA on Matrix Market input

This example performs CA on the sparse MM input (Figure 3). `sparse_ca` reads (row, col, value) triplets in chunks and applies the same implicit operator without ever materializing dense intermediates of size $N \times M$.

```
out_sparse_ca = sparse_ca(input=joinpath(tmp, "Data.mtx.zst"),
                          dim=3, noversamples=5, niter=3, chunksize=100)
```

```
subplots(out_sparse_ca, group)
```




Figure 3: Output of `sparse_ca` against binarized MM format.

For more details, including the BinCOO variant for binary contingency tables, see the README.md of `OnlineCA.jl` at <https://github.com/chiba-ai-med/OnlineCA.jl>.

Related work

There are various implementations of CA (Esteves, 2017; Halford, 2024; Lê et al., 2008; Nenadić & Greenacre, 2007; Oksanen & others, 2024) and some of them support sparse data formats (DeBruine, 2023; Hsu & Culhane, 2023), but `OnlineCA.jl` is the only tool that supports both OOC computation and language-agnostic sparse formats (MM and BinCOO), enabling seamless integration with external data pipelines.

Function Name	Language	OOC	Sparse Format
<code>MASS::corresp</code>	R	No	-
<code>ca::ca</code>	R	No	-
<code>FactoMineR::CA</code>	R	No	-
<code>vegan::cca</code>	R	No	-
<code>prince.CA</code>	Python	No	-
<code>mca</code>	Python	No	-
<code>corral</code>	R	No	dgCMatrix
<code>singlet</code>	R	No	dgCMatrix
<code>OnlineCA.jl</code>	Julia	Yes	MM/BinCOO

References

- Braak, C. J. F. ter. (1986). Canonical correspondence analysis: A new eigenvector technique for multivariate direct gradient analysis. *Ecology*, 67(5), 1167–1179. <https://doi.org/10.2307/1938672>
- DeBruine, Z. J. (2023). *singlet: Non-negative matrix factorization and correspondence analysis for single-cell genomics*. <https://github.com/zdebruine/singlet>.
- Esteves, E. (2017). *mca: Multiple correspondence analysis in Python*. <https://github.com/esafak/mca>.

- Greenacre, M. (2017). *Correspondence analysis in practice* (3rd ed.). Chapman; Hall/CRC. <https://doi.org/10.1201/9781315369983>
- Halford, M. (2024). *prince: Multivariate exploratory data analysis in Python*. <https://github.com/MaxHalford/prince>.
- Halko, N., Martinsson, P. G., Shkolnisky, Y., & Tygert, M. (2011). An algorithm for the principal component analysis of large data sets. *SIAM Journal on Scientific Computing*, 33(5), 2580–2594. <https://doi.org/10.1137/100804139>
- Halko, N., Martinsson, P. G., & Tropp, J. A. (2011). Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM Review*, 53(2), 217–288. <https://doi.org/10.1137/090771806>
- Hill, M. O. (1974). Correspondence analysis: A neglected multivariate method. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 23(3), 340–354. <https://doi.org/10.2307/2347127>
- Hsu, L. L., & Culhane, A. C. (2023). Corral: Single-cell RNA-seq dimension reduction, batch integration, and visualization with correspondence analysis. *Bioinformatics*, 39(2), btad027. <https://doi.org/10.1093/bioinformatics/btad027>
- Lê, S., Josse, J., & Husson, F. (2008). FactoMineR: An R package for multivariate analysis. *Journal of Statistical Software*, 25(1), 1–18. <https://doi.org/10.18637/jss.v025.i01>
- Nenadić, O., & Greenacre, M. (2007). Correspondence analysis in R, with two- and three-dimensional graphics: The ca package. *Journal of Statistical Software*, 20(3), 1–13. <https://doi.org/10.18637/jss.v020.i03>
- Oksanen, J., & others. (2024). *vegan: Community ecology package*. <https://CRAN.R-project.org/package=vegan>. <https://doi.org/10.32614/CRAN.package.vegan>
- Tsuyuzaki, K. (2025). *OnlineNMF.jl: A Julia package for out-of-core and sparse non-negative matrix factorization*. <https://github.com/rikenbit/OnlineNMF.jl>. <https://doi.org/10.5281/zenodo.8199413>
- Tsuyuzaki, K. (2026). OnlinePCA.jl: A Julia package for out-of-core and sparse principal component analysis. *Journal of Open Source Software*, 11(117), 9343. <https://doi.org/10.21105/joss.09343>
- Tsuyuzaki, K., Sato, H., Sato, K., & Nikaido, I. (2020). Benchmarking principal component analysis for large-scale single-cell RNA-sequencing. *Genome Biology*, 21(1). <https://doi.org/10.1186/s13059-019-1900-3>